

Statistical Methods for Insurance: Compiling data for problem solving

Di Cook & Souhaib Ben Taieb, Econometrics and Business Statistics, Monash University
W8.C1

Overview of this class

- What is `tidy` data? Why do you want tidy data? Getting your data into tidy form using `tidyr`.
- Wrangling verbs: `filter`, `arrange`, `select`, `mutate`, `summarise`, with `dplyr`
- Date and time with `lubridate`

Terminology

1. Cases, records, individuals, subjects, experimental units, example, instance: things we are collecting information about
2. Variables, attributes, fields, features: what we are measuring on each record/case/.../instance

Generally we think of cases being on the rows, and variables being in the columns of a table. This is a basic data structure. BUT data often is given to us in many other shapes than this. Getting into a tidy shape will allow you to efficiently use it for modeling.

Example 1

Inst	AvNumPubs	AvNumCits	PctCompletion
ARIZONA STATE UNIVERSITY	0.90	1.57	31.7
AUBURN UNIVERSITY	0.79	0.64	44.4
BOSTON COLLEGE	0.51	1.03	46.8
BOSTON UNIVERSITY	0.49	2.66	34.2

- Cases: _____
- Variables: _____

Example 2

V1	V2	V3	V4	V5	V9	V13	V17	V21	V25	V29	V33	V37	V41	V45	V49	V53	V57
ASN00086282	1970	7	TMAX	141	124	113	123	148	149	139	153	123	108	119	112	126	112
ASN00086282	1970	7	TMIN	80	63	36	57	69	47	84	78	49	42	48	56	51	36
ASN00086282	1970	7	PRCP	3	30	0	0	36	3	0	0	10	23	3	0	5	0
ASN00086282	1970	8	TMAX	145	128	150	122	109	112	116	142	166	127	117	127	159	143

- Cases: _____
- Variables: _____

Example 3

Here are the column headers ...

```
#> [1] "iso2"    "year"    "m_04"    "m_514"   "m_014"   "m_1524"  "m_2534"
#> [8] "m_3544"  "m_4554"  "m_5564"  "m_65"    "m_u"     "f_04"    "f_514"
#> [15] "f_014"   "f_1524"  "f_2534"  "f_3544"  "f_4554"  "f_5564"  "f_65"
#> [22] "f_u"
```

- Cases: _____
- Variables: _____

Example 4

We'll commonly find these data on web sites:

religion	<\$10k	\$10-20k	\$20-30k	\$30-40k
Agnostic	27	34	60	81
Atheist	12	27	37	52
Buddhist	27	21	30	34
Catholic	418	617	732	670
Don't know/refused	15	14	15	11

- Cases: _____
- Variables: _____

Example 5

10 week sensory experiment, 12 individuals assessed taste of french fries on several scales (how potato-y, buttery, grassy, rancid, paint-y do they taste?), fried in one of 3 different oils, replicated twice. First few rows:



time	treatment	subject	rep	potato	buttery	grassy	rancid	painty
1	1	3	1	2.9	0.0	0.0	0.0	5.5
1	1	3	2	14.0	0.0	0.0	1.1	0.0
1	1	10	1	11.0	6.4	0.0	0.0	0.0
1	1	10	2	9.9	5.9	2.9	2.2	0.0

What do you like to know?

12 subjects



Three oils,
two batches



Five scales



RANCID



For 10 weeks

S	M	T	W	T	F	S
28	29	1	2	3	4	5
6	7	8	9	10	11	12
13	14	15	16	17	18	19
20	21	22	23	24	25	26
27	28	29	30	31	1	2
27	28	29	30	31	1	2
3	4	5	6	7	8	9
10	11	12	13	14	15	16
17	18	19	20	21	22	23
24	25	26	27	28	29	30
1	2	3	4	5	6	7
8	9	10	11	12	13	14
15	16	17	18	19	20	21
22	23	24	25	26	27	28
29	30	31	1	2	3	4
..

Messy Data Patterns

There are various features of messy data that one can observe in practice. Here are some of the more commonly observed patterns.

- Column headers are values, not variable names
- Variables are stored in both rows and columns, contingency table format
- Information stored in multiple tables
- Dates in many different formats
- Not easy to analyse

What is Tidy Data?

- Each observation forms a row
- Each variable forms a column
- Contained in a single table
- Long form makes it easier to reshape in many different ways
- Wide form is common for analysis/modeling

Description by Hadley Wickham

Tidy data = lego

<http://www.flickr.com/photos/wwworks/2473052504>

Description by Hadley Wickham

Messy data = play mobile



<https://www.flickr.com/photos/kafka4prez/57282282>

Tidy vs Messy

- Tidy data facilitates analysis in many different ways, answering multiple questions, applying methods to new data or other problems
- Messy data may work for one particular problem but is not generalisable

Tidy Verbs

- **gather**: specify the **keys** (identifiers) and the **values** (measures) to make long form (used to be called melting)
- **spread**: variables in columns (used to be called casting)
- **nest/unnest**: working with lists
- **separate/unite**: split and combine columns

French fries example


	time	treatment	subject	rep	potato	buttery	grassy	rancid	painty
61	1	1	3	1	2.9	0.0	0.0	0.0	5.5
25	1	1	3	2	14.0	0.0	0.0	1.1	0.0
62	1	1	10	1	11.0	6.4	0.0	0.0	0.0
26	1	1	10	2	9.9	5.9	2.9	2.2	0.0
63	1	1	15	1	1.2	0.1	0.0	1.1	5.1
27	1	1	15	2	8.8	3.0	3.6	1.5	2.3

This format is not ideal for data analysis

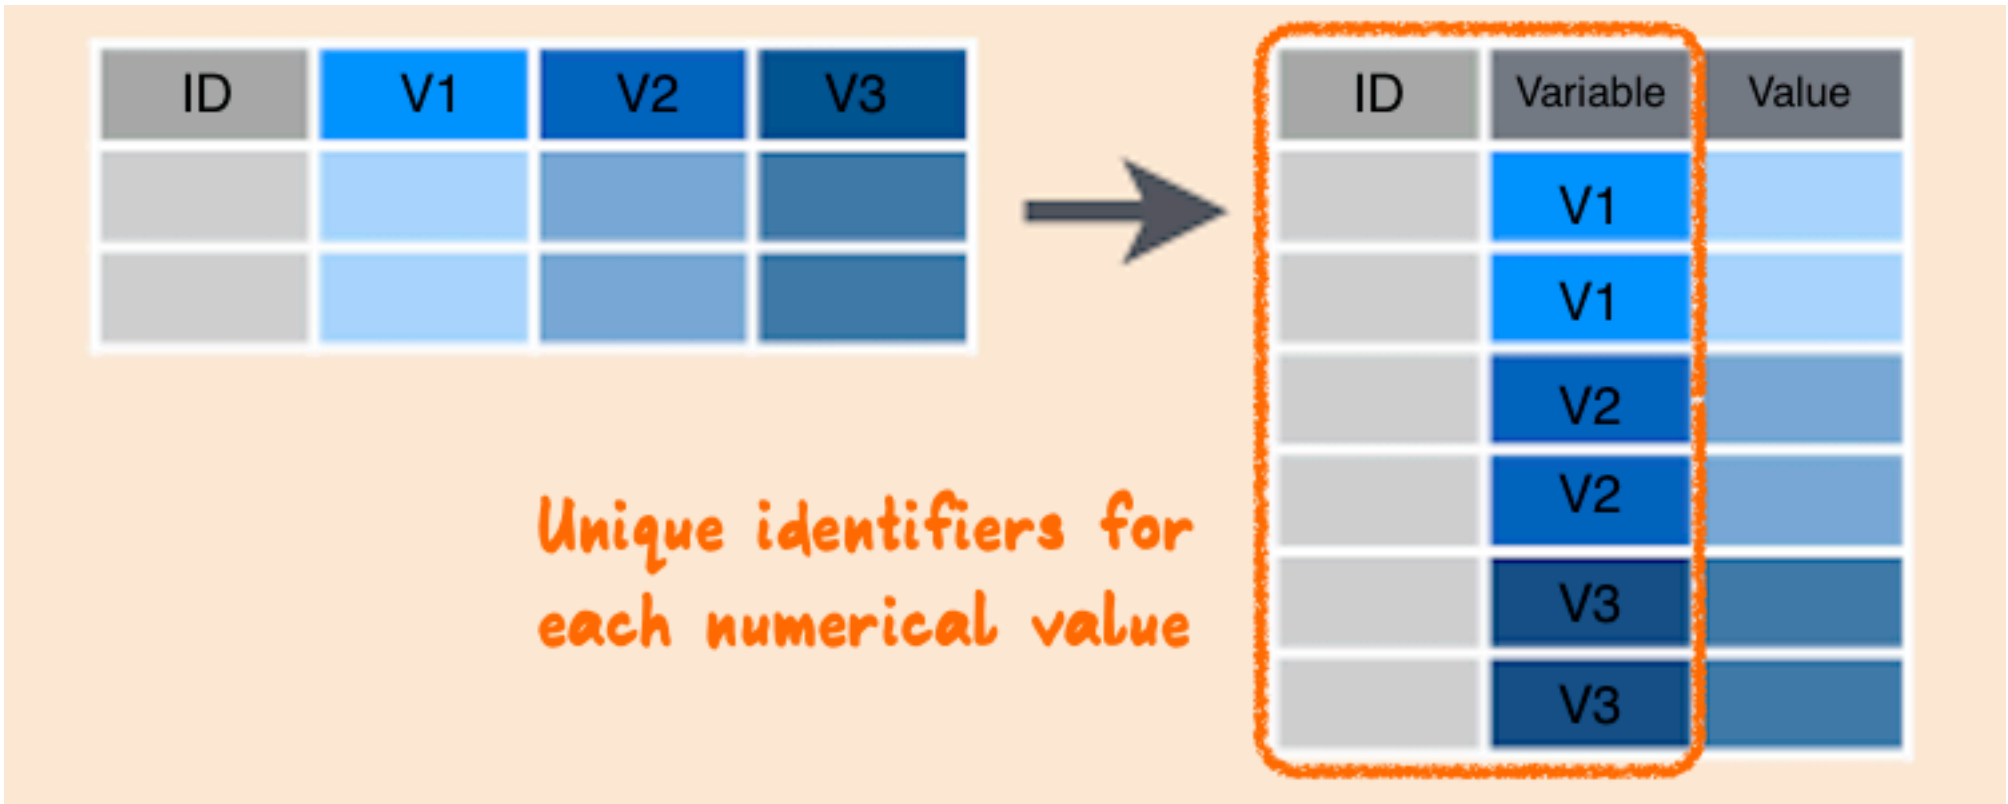
What code would be needed to plot each of the ratings over time as a different color?

```
library(ggplot2)
french_sub <- french_fries[french_fries$time == 10,]
ggplot(data = french_sub) +
  geom_boxplot(aes(x="1_potato", y=potato), fill = I("red")) +
  geom_boxplot(aes(x = "2_buttery", y = buttery), fill = I("orange")) +
  geom_boxplot(aes(x = "3_grassy", y = grassy), fill = I("yellow")) +
  geom_boxplot(aes(x = "4_rancid", y = rancid), fill = I("green")) +
  geom_boxplot(aes(x = "5_painty", y = painty), fill = I("blue")) +
  xlab("variable") + ylab("rating")
```

The plot



Wide to long



Gathering

- When gathering, you need to specify the keys (identifiers) and the values (measures).
- Keys/Identifiers:
 - Identify a record (must be unique)
 - Example: Indices on an random variable
 - Fixed by design of experiment (known in advance)
 - May be single or composite (may have one or more variables)
- Values/Measures:
 - Collected during the experiment (not known in advance)
 - Usually numeric quantities

Gathering the French Fries Data

```
ff_long <- gather(french_fries, key = variable,
  value = rating, potato:painty)


head(ff_long)
```

#>	time	treatment	subject	rep	variable	rating
#> 1	1	1	3	1	potato	2.9
#> 2	1	1	3	2	potato	14.0
#> 3	1	1	10	1	potato	11.0
#> 4	1	1	10	2	potato	9.9
#> 5	1	1	15	1	potato	1.2
#> 6	1	1	15	2	potato	8.8

Let's re-write the code for our Plot

```
ff_long_sub <- ff_long[  
  french_fries_long$time == 10,]  
  
ggplot(data = ff_long_sub,  
  aes(x=variable, y=rating, fill = variable)) +  
  geom_boxplot()
```

And plot it



Long to Wide

In certain applications, we may wish to take a long dataset and convert it to a wide dataset (Perhaps displaying in a table).

```
#>   time treatment subject rep variable rating
#> 1     1           1       3    1  potato    2.9
#> 2     1           1       3    2  potato   14.0
#> 3     1           1      10    1  potato   11.0
#> 4     1           1      10    2  potato    9.9
#> 5     1           1      15    1  potato    1.2
#> 6     1           1      15    2  potato    8.8
```


Spread

We use the spread function from tidyr to do this:

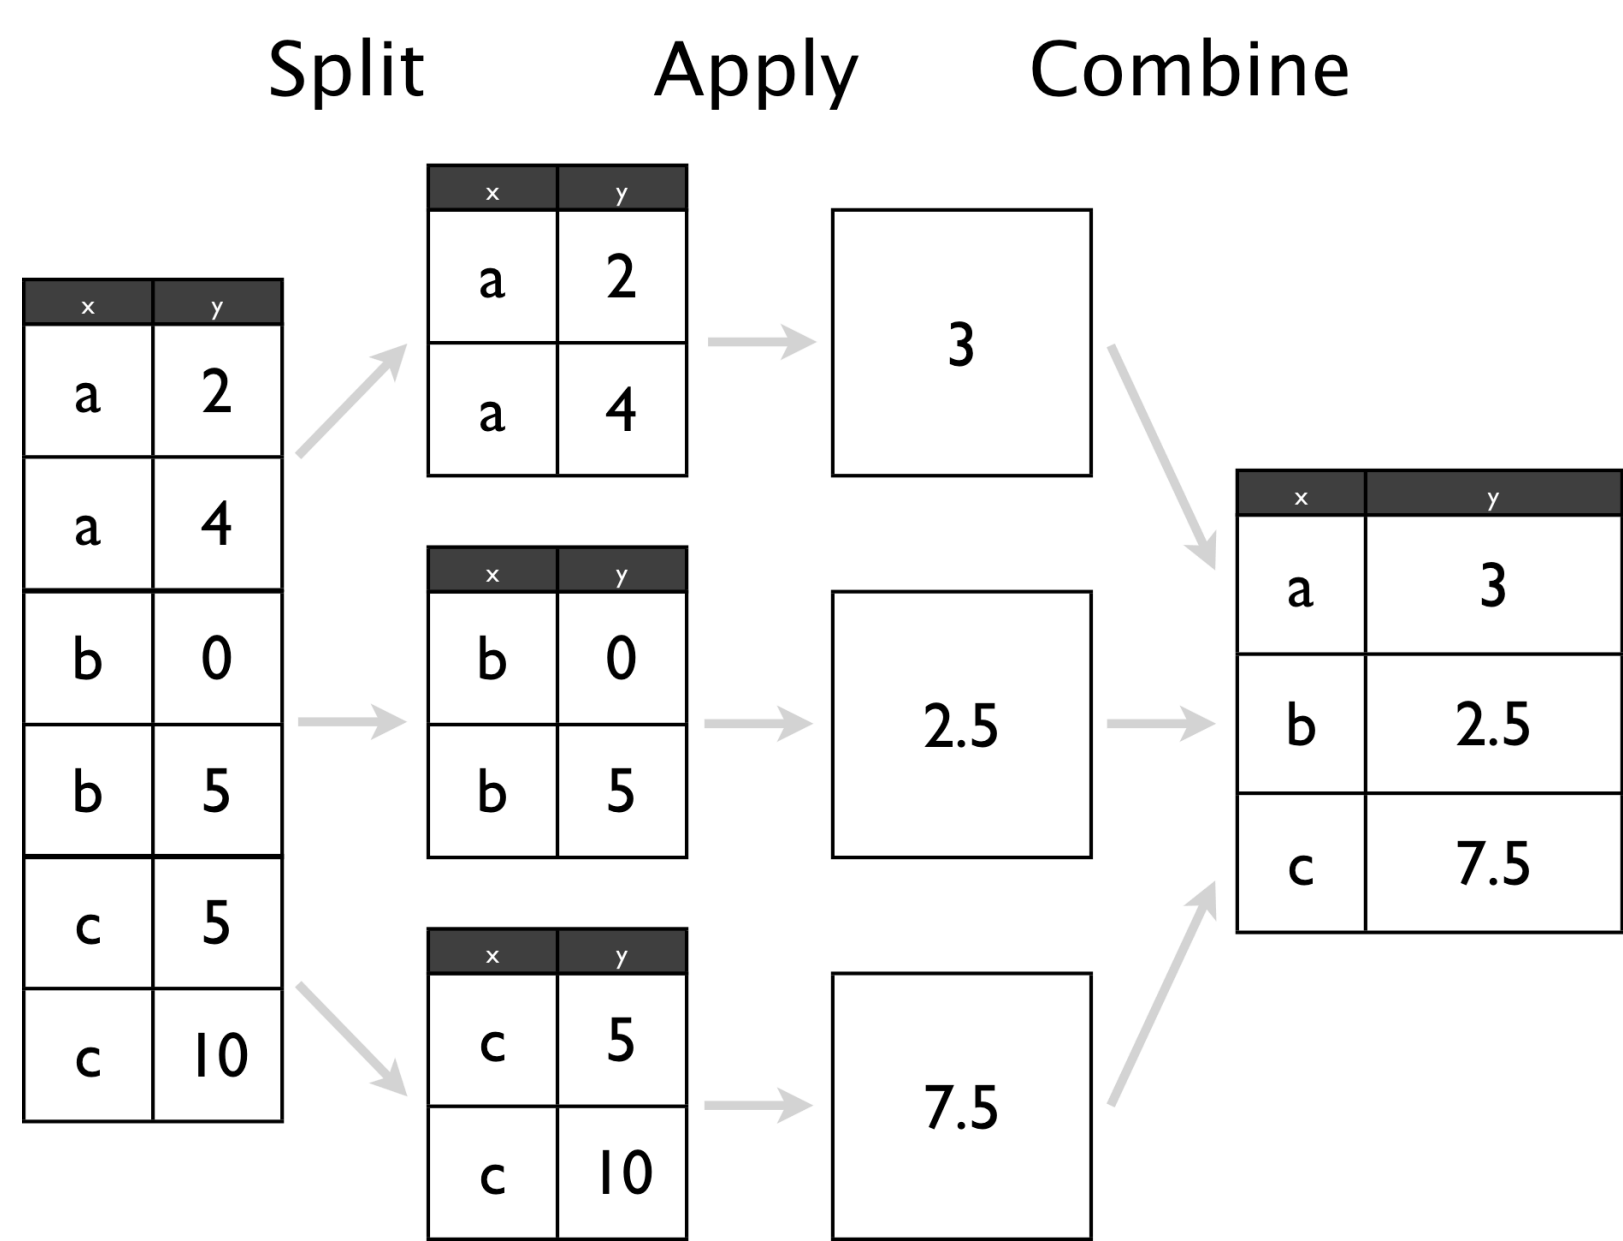
```
ff_wide <- spread(ff_long,
  key = variable, value = rating)
head(ff_wide)
```

```
#>   time treatment subject rep buttery grassy painty potato rancid
#> 1     1           1       3     1     0.0    0.0    5.5     2.9     0.0
#> 2     1           1       3     2     0.0    0.0    0.0    14.0     1.1
#> 3     1           1      10     1     6.4    0.0    0.0    11.0     0.0
#> 4     1           1      10     2     5.9    2.9    0.0     9.9     2.2
#> 5     1           1      15     1     0.1    0.0    5.1     1.2     1.1
#> 6     1           1      15     2     3.0    3.6    2.3     8.8     1.5
```

The Split-Apply-Combine Approach

- Split a dataset into many smaller sub-datasets
- Apply some function to each sub-dataset to compute a result
- Combine the results of the function calls into a one dataset

The Split-Apply-Combine Approach



Split-Apply-Combine in dplyr

```
library(dplyr)
ff_summary <- group_by(ff_long, variable) %>% # SPLIT
  summarise(
    m = mean(rating, na.rm = TRUE),
    s=sd(rating, na.rm=TRUE)) # APPLY + COMBINE
ff_summary
```

```
#> # A tibble: 5 x 3
#>   variable      m      s
#>   <chr>    <dbl>  <dbl>
#> 1 buttery 1.8236994 2.409758
#> 2 grassy  0.6641727 1.320574
#> 3 painty  2.5217579 3.393717
#> 4 potato  6.9525180 3.584403
#> 5 rancid   3.8522302 3.781815
```

Pipes

- Pipes historically enable data analysis pipelines
- Pipes allow the code to be read like a sequence of operations
- dplyr allows us to chain together these data analysis tasks using the `%>%` (pipe) operator
- `x %>% f(y)` is shorthand for `f(x, y)`
- Example:

```
student2012.sub <- readRDS("../data/student_sub.rds")
student2012.sub %>%
  group_by(CNT) %>%
  tally()
#> # A tibble: 43 x 2
#>   CNT      n
#>   <chr> <int>
#> 1  ARE 11500
#> 2  AUS 14481
#> 3  AUT  4755
#> 4  BEL  8597
#> 5  BGR  5282
#> 6  BRA  5506
#> 7  CAN 21544
#> 8  CHL  6856
#> 9  COL  9073
#> 10 CZE  5327
#> # ... with 33 more rows
```

dplyr verbs

There are five primary dplyr **verbs**, representing distinct data analysis tasks:

- **Filter**: Remove the rows of a data frame, producing subsets
- **Arrange**: Reorder the rows of a data frame
- **Select**: Select particular columns of a data frame
- **Mutate**: Add new columns that are functions of existing columns
- **Summarise**: Create collapsed summaries of a data frame

Filter

```
french_fries %>%  
  filter(subject == 3, time == 1)
```

#>	time	treatment	subject	rep	potato	buttery	grassy	rancid	painty
#> 1	1		1	3	1	2.9	0.0	0.0	5.5
#> 2	1		1	3	2	14.0	0.0	0.0	1.1
#> 3	1		2	3	1	13.9	0.0	0.0	3.9
#> 4	1		2	3	2	13.4	0.1	0.0	1.5
#> 5	1		3	3	1	14.1	0.0	0.0	1.1
#> 6	1		3	3	2	9.5	0.0	0.6	2.8

Arrange

```
french_fries %>%  
  arrange(desc(rancid)) %>%  
  head
```

#>	time	treatment	subject	rep	potato	buttery	grassy	rancid	painty
#> 1	9	2	51	1	7.3	2.3	0	14.9	0.1
#> 2	10	1	86	2	0.7	0.0	0	14.3	13.1
#> 3	5	2	63	1	4.4	0.0	0	13.8	0.6
#> 4	9	2	63	1	1.8	0.0	0	13.7	12.3
#> 5	5	2	19	2	5.5	4.7	0	13.4	4.6
#> 6	4	3	63	1	5.6	0.0	0	13.3	4.4

Select

```
french_fries %>%
  select(time, treatment, subject, rep, potato) %>%
  head
```

```
#>      time treatment subject rep potato
#> 61      1          1       3    1    2.9
#> 25      1          1       3    2   14.0
#> 62      1          1      10    1   11.0
#> 26      1          1      10    2    9.9
#> 63      1          1      15    1    1.2
#> 27      1          1      15    2    8.8
```

Mutate

```
french_fries %>%
  mutate(yucky = grassy+rancid+painty) %>%
  head
```

```
#>   time treatment subject rep potato buttery grassy rancid painty yucky
#> 1     1           1         3     1    2.9     0.0    0.0    0.0    5.5    5.5
#> 2     1           1         3     2   14.0     0.0    0.0    1.1    0.0    1.1
#> 3     1           1        10     1   11.0     6.4    0.0    0.0    0.0    0.0
#> 4     1           1        10     2    9.9     5.9    2.9    2.2    0.0    5.1
#> 5     1           1        15     1    1.2     0.1    0.0    1.1    5.1    6.2
#> 6     1           1        15     2    8.8     3.0    3.6    1.5    2.3    7.4
```

Summarise

```
french_fries %>%
  group_by(time, treatment) %>%
  summarise(mean_rancid = mean(rancid),
            sd_rancid = sd(rancid))
```

```
#> Source: local data frame [30 x 4]
#> Groups: time [?]
```

```
#>
```

#>	time	treatment	mean_rancid	sd_rancid
#>	<fctr>	<fctr>	<dbl>	<dbl>
#> 1	1	1	2.758333	3.212870
#> 2	1	2	1.716667	2.714801
#> 3	1	3	2.600000	3.202037
#> 4	2	1	3.900000	4.374730
#> 5	2	2	2.141667	3.117540
#> 6	2	3	2.495833	3.378767
#> 7	3	1	4.650000	3.933358
#> 8	3	2	2.895833	3.773532
#> 9	3	3	3.600000	3.592867
#> 10	4	1	2.079167	2.394737

```
#> # ... with 20 more rows
```

Dates and Times

- Dates are deceptively hard to work with
- 02/05/2012. Is it February 5th, or May 2nd?
- Time zones
- Different starting times of stock markets, airplane departure and arrival

Basic Lubridate Use

```
library(lubridate)

now()
#> [1] "2016-09-14 19:52:52 AEST"
now(tz = "America/Chicago")
#> [1] "2016-09-14 04:52:52 CDT"
today()
#> [1] "2016-09-14"
now() + hours(4)
#> [1] "2016-09-14 23:52:52 AEST"
today() - days(2)
#> [1] "2016-09-12"
ymd("2013-05-14")
#> [1] "2013-05-14"
mdy("05/14/2013")
#> [1] "2013-05-14"
dmy("14052013")
#> [1] "2013-05-14"
```

Dates example: Oscars date of birth

```
oscars <- read_csv("../data/oscars.csv")
oscars <- oscars %>% mutate(DOB = mdy(DOB))
head(oscars$DOB)
#> [1] "1895-09-30" "1884-07-23" "1894-04-23" "2006-10-06" "1886-02-02"
#> [6] "1892-04-08"
summary(oscars$DOB)
#>      Min.      1st Qu.      Median      Mean      3rd Qu.
#> "1868-04-10" "1934-09-18" "1957-06-23" "1962-05-21" "2008-04-05"
#>      Max.
#> "2029-12-13"
```

Calculating on dates

- You should never ask a woman her age, but ... really!


```
oscars <- oscars %>% mutate(year=year(DOB))
summary(oscars$year)
#>      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
#>      1868   1934   1957   1962   2008   2029
oscars %>% filter(year == "2029") %>%
  select(Name, Sex, DOB)
#> # A tibble: 4 x 3
#>       Name      Sex      DOB
#>   <chr>   <chr>   <date>
#> 1 Audrey Hepburn Female 2029-05-04
#> 2 Grace Kelly Female 2029-11-12
#> 3 Miyoshi Umeki Female 2029-04-03
#> 4 Christopher Plummer Male 2029-12-13
```

Months

```
oscars <- oscars %>% mutate(month=month(DOB, label = TRUE, abbr = TRUE))
table(oscars$month)
#>
#> Jan Feb Mar Apr May Jun Jul Aug Sep Oct Nov Dec
#>  32  35  37  53  35  27  37  33  31  31  30  42
```



Now plot it

```
ggplot(data=oscars, aes(month)) + geom_bar()
```



Should you be born in April?

```
df <- data.frame(m=sample(1:12, 423, replace=TRUE))  
df$m2 <- factor(df$m, levels=1:12,  
  labels=month.abb)  
ggplot(data=df, aes(x=m2)) + geom_bar()
```



Resources

- [Tidy data](#)
- [Split-apply-combine](#)
- [RStudio cheat sheets](#)
- [Working with dates and times](#)
- [R for Data Science](#)

Share and share alike

This work is licensed under the Creative Commons Attribution-Noncommercial 3.0 United States License. To view a copy of this license, visit <http://creativecommons.org/licenses/by-nc/3.0/us/> or send a letter to Creative Commons, 171 Second Street, Suite 300, San Francisco, California, 94105, USA.